

## Homework 1 analysis report

### 1. Description of datasets :

#### a. 20NewsGroups data :

One of the popular dataset for text mining. Has documents of 20 news groups. Each group has about 1000 documents within them. Some of these new groups are closely related while some of them totally unrelated. Below table summarizes the related groups:

comp.graphics comp.os.ms-windows.misc comp.sys.ibm.pc.hardware comp.sys.mac.hardware comp.windows.x	rec.autos rec.motorcycles rec.sport.baseball rec.sport.hockey	sci.crypt sci.electronics sci.med sci.space
misc.forsale	talk.politics.misc talk.politics.guns talk.politics.mideast	talk.religion.misc alt.atheism soc.religion.christian

I took the [20news-19997.tar.gz](#)

I took the original dataset which has contains duplicate documents and header in each documents.

Upon, looking into the documents in the groups I found that header has fields like subject, date, organization etc. Since each document has these terms, we have to remove these words during preprocessing as these does not give any information if the document itself.

#### b. Yelp dataset :

This dataset is recent real time dataset collected from the yelp user reviews. This has 6 parts, namely, business, review, user, checkin and tip(each of which are json files).

Upon examining, reviews can be grouped as positive, negative, sarcastic, suggestive etc We were provided with the csv files created from the jsons of reviews and business for simplicity and I am using these csv files.

Business csv contains business ids and categories files from the business json

Reviews csv contains business ids and reviews from the review json

Topics	No. of Documents	No. of Unique terms	No. of unique terms after removing stop words	No. of unique terms after Stemming
alt.atheism	480	10879	10808	7427
comp.graphics	584	12076	12010	9400
comp.os.ms-windows.misc	591	24271	24206	22401
comp.sys.ibm.pc.hardware	590	9234	9168	7305
comp.sys.mac.hardware	578	8557	8490	6560
comp.windows.x	593	13265	13199	10658
misc.forsale	585	9734	9666	7875
rec.autos	594	10475	10405	7844
rec.motorcycles	598	10490	10422	7884
rec.sport.baseball	597	9225	9157	7122
rec.sport.hockey	600	11399	11329	8833
sci.crypt	595	13545	13475	9610
sci.electronics	591	10219	10153	7638
sci.med	594	14413	14342	10566
sci.space	593	13804	13735	9921

soc.religion.christian	599	12800	12728	8697
talk.politics.guns	546	12905	12835	9152
talk.politics.mideast	564	15458	15387	10838
talk.politics.misc	465	12705	12634	8712
talk.religion.misc	377	10815	10746	7616

## 2. Experiments :

### a. Data Processing :

For 20NG dataset I am Using below groups:

rec.sport.hockey

rec.sport.baseball

sci.space

Among these rec.sport.hockey and rec.sport.baseball and similar to each other

And sci.space is dissimilar from both rec.sport.hockey and rec.sport.baseball

For Yelp dataset I am choosing Infer Categories subproblem and using data from below categories:

American

Asian

Thai

Sports

Among these American, Asian and Thai are similar as these are types of restaurants

Sports is dissimilar from all three.

#loading selected data from 20NG data set from the folders

```
data2<-
```

```
c("20_newsgroups\\rec.sport.hockey","20_newsgroups\\rec.sport.baseball","20_newsgroups\\sci.space") #Read the text files for 3 groups
```

```
news <- Corpus(DirSource(data2, recursive=TRUE),readerControl = list(reader=readPlain))# Load the data as a corpus
```

```
#loading selected Yelp data from the business.csv and reviews.csv files provided too us
business.category <- read.csv (file="yelp_academic_dataset_business.csv", header=TRUE)
```

```
business.reviews <- read.csv (file = "yelp_academic_dataset_review.csv", header=FALSE ,nrows
= 5000000,sep="," )
```

```
#Selecting Asian , American, Thai and Sports categories
business.seletedCategories <-
business.category[grepl("American",business.category$categories)
                  | grepl("Asian",business.category$categories)
                  | grepl("Thai",business.category$categories)
                  | grepl("Sports",business.category$categories),]
```

```
#removing b' at the beginning and ' at the end of business_id in business.csv
business.category$business_id <- gsub("^b", "", business.category$business_id)
business.category$business_id <- gsub("'", "", business.category$business_id)
```

```
#merging both categories and reviews in the common business_id field
filtered <- merge ( x = business.seletedCategories , y = business.reviews, by.x=c("business_id"),
by.y=c("business_id"))
```

```
#collapsing the filtered data on the categories, so that reviews for similar categories are placed
together
collapsed <- aggregate(text ~categories, data = filtered, toString)
```

```
t <- data.frame(collapsed$text, stringsAsFactors = FALSE)
corpus <- Corpus(DataframeSource(t))
corpus <- tm_map(corpus, stripWhitespace)
corpus <- tm_map(corpus, content_transformer(tolower))
corpus <- tm_map(corpus, removeWords, stopwords("english"))
corpus <- tm_map(corpus, stemDocument)
corpus <- tm_map(corpus, removeNumbers)
corpus <- tm_map(corpus, removePunctuation)
corpus <- tm_map(corpus,content_transformer(PlainTextDocument))
```

Creating DocumentTermMatrix :

```
dtm <- DocumentTermMatrix(news,control=list(wordLengths=c(4,Inf)))#considering only terms
whose length is greater thann 4
```

Creating TermDocumentMatrix :

```
tdm <- TermDocumentMatrix(news, control=list(wordLengths=c(4,Inf)))
```

Document-term matrix is an  $M \times N$  matrix, where  $M$  is number of documents and  $N$  is number of words in the entire corpus, and hold number of times a word is occurring in a particular document.

Term-document matrix is transpose of this matrix.

Below is a snapshot of DTM created from the corpus(Showing for 20Groups Data):

```
> dtm #document matrix for NewsGroups I have selected
<<DocumentTermMatrix (documents: 3000, terms: 84623)>>
Non-/sparse entries: 424546/253444454
Sparsity           : 100%
Maximal term length: 295
weighting          : term frequency (tf)
```

Top 10 words from the Term-document matrix :

	word freq
that	that 7055
have	have 3868
this	this 3726
with	with 3322
from:	from: 3041
date:	date: 3019
subject:	subject: 3015
newsgroups:	newsgroups: 3011
message-id:	message-id: 3009
path:	path: 3003

In the preprocessing steps we remove all unwanted data which do not give any information and clean the data i.e, converting the text too lower case, remove punctuations, stopwords, remove some keywords repeating in all the documents, do stemming to reduce the words to their root forms, remove white spaces etc.

Below are the snapshots of DTM are processing(before stemming):

```
> dtm.preprocessed.notstemmed <- DocumentTermMatrix(news,control=list(wordLengths=c(4,Inf)))
> dtm.preprocessed.notstemmed
<<DocumentTermMatrix (documents: 3000, terms: 36649)>>
Non-/sparse entries: 278371/109668629
Sparsity           : 100%
Maximal term length: 234
weighting          : term frequency (tf)
```

Here we can clearly see the difference in the number of terms. Preprocessing (removing stopwords, common words (headers for example), punctuations, whitespaces ) is an important step and has a significant impact on the data that we are operating on.

Top 10 words from TDM after preprocessing and before stemming :

	word	freq
will	will	2257
game	game	1810
space	space	1765
university	university	1599
team	team	1584
like	like	1472
just	just	1468
think	think	1349
year	year	1344
first	first	1209

Below is the code snippet from my solution which does preprocessing as mentioned above :

```
#Preprocessing the data for 20NG dataset
skipWords <- function(x) removeWords(x,
c("xref","path","from","newsgroups","subject","message-id","messageid","date","article-
id","article","references","organization","nntp-posting-host","nntppostinghost","lines",
"expires","reply-to","followup-to","distribution","keywords","summary","sender","writes",
stopwords("english")))
funcs <- list(removePunctuation, stripWhitespace, removeNumbers,
skipWords,content_transformer(tolower))
news <- tm_map(news, FUN = tm_reduce, tmFuns = funcs)
news <- tm_map(news,content_transformer( PlainTextDocument))

#preprocessing for yelp dataset
t <- data.frame(collapsed$text, stringsAsFactors = FALSE)
corpus <- Corpus(DataframeSource(t))
corpus <- tm_map(corpus, stripWhitespace)
corpus <- tm_map(corpus, content_transformer(tolower))
corpus <- tm_map(corpus, removeWords, stopwords("english"))
corpus <- tm_map(corpus, stemDocument)
corpus <- tm_map(corpus, removeNumbers)
corpus <- tm_map(corpus, removePunctuation)
corpus <- tm_map(corpus,content_transformer(PlainTextDocument))
```

Below are the snapshots of DTM affter Stemming) and the word cloud for the same:

Inf)))

words into their root forms , for ex. Education will be converted to educate.

Top 10 words from TDM after stemming :

game	game	2843
will	will	2325
team	team	2117
year	year	2112
space	space	1771
univers	univers	1722
like	like	1678
play	play	1507
think	think	1499
time	time	1479



Typically even after preprocessing and stemming DTM will be very sparse (has 0 or some low value) as many words occur only in few documents(We can observe the same from the above snapshot of my DTM for 20Groups data ).

So we do remove sparse terms to drop those very less frequent words by doing below step :

```
dtm.sparse = removeSparseTerms(dtm, 0.75)# we are removing terms which are which are not occurring in 97% of the documents atleast.
```

Many common words like articles, conjunctions occur more frequently in documents than other words. Therefore, using counts can be deceiving. To overcome these problem and punish these words for occurring in all the documents we do term frequency-inverse document frequency weight instead of counts as shown below :

```
tdm.tf.idf.stemmed <- weightTfIdf(tdm.stemmed)
```

After all these steps we have our data ready for experimenting and finding the hidden information.

### 3. Experiments :

Here I am presenting the results I got when by plotting the clusters directly from the data, using Latent Semantic Analysis(A Dimensionality Reduction Technique) with Singular Value Decomposition and Latent Dirichlet Allocation(another DR technique).

#### Latent Semantic Analysis with SVD:

This technique decomposes our DTM into 3 matrices and uncovers the space represented by it. This method find hidden concepts in the data.

$$A = U d V^{\text{transpose}}$$

Concepts are represented by the singular vectors

Each document and word are reprinted in terms of the concept vectors, i.e. in terms of projection on the singular vector representing ■

Then we reduce all three matrices into k dimension, by selecting only k largest eigen values in the diagonal matrix d.



Below is the R code for the same :

```
svd <-svd(dtm.tf.idf.sparsed)
D <- diag(svd$d)
numberOfTerms = 200#,100,200
diagD <- D[1:numberOfTerms,1:numberOfTerms]
documentVector<-svd$u[,1:numberOfTerms] %*% diagD #creating the document vector
wordVector<-diagD %*%t(svd$v[,1:numberOfTerms]) #creating the word vector
```

Since U contains the eigen values of document-document matrix, this when multiplied with the diagonal matrix containing the eigen values, gives the importance of document

Since V contains the eigen values of document-document matrix, this when multiplied with the diagonal matrix containing the eigen values, gives the importance of term

## K-means Clustering :

Is a prototype based, partition based clustering technique that attempts to find the user-specified number of clusters in the data, which are represented by their centroids.

We must have an objective function to measure the quality of clustering. In this case we use Sum of Squares Error (SSE), which is also known as SCATTER. In SSE we calculate the error of each data point, i.e. its Euclidean distance to the closest centroid and then compute the total sum of squared error by summing over for all datapoints in all clusters.

Plotting k-means clusters in R:

```
m <- as.matrix(dtm.tf.idf.stemmed)
rownames(m) <- 1:nrow(m)
norm_eucl<- function(m)m/apply(m,MARGIN=1, FUN=function(x)sum(x^2)^.5)
m_norm<-norm_eucl(m)
cl<-kmeans(m_norm,2)
table(cl$cluster)
plot(prcomp(m_norm)$x, col=cl$cl)
```

## Latent Dirichlet Allocation :

LDA is a generative probabilistic model of the corpus.

LDA models how the data was generated in order to categorize a signal. LDA asks the question to itself : based on my generation assumption, which categories are more likely to generate this signal.

Basic idea on which LDA operates is that : documents are generated as random mixtures over latent topics, where each topic is characterized by a distribution over words.

LDA involves 3-levels, and notably the topic node is sampled repeatedly within the document. Under this model documents can be associated with multiple topics.

My Implementation in R :

```
lda <- LDA(dtm,300,method="Gibbs")
```

```
lda.topics.300 <- as.matrix(topics(lda))
```

```
head(lda.topics.300,20)
```

```
lda.terms <- as.matrix(terms(lda,))
```

```
head(lda.terms,20)
```

```
topicProbabilities <- as.data.frame(lda@gamma)
```

```
m <- as.matrix(topicProbabilities)
```

```
rownames(m) <- 1:nrow(m)
```

```
norm_eucl <- function(m) m/apply(m, MARGIN=1, FUN=function(x) sum(x^2)^.5)
```

```
m_norm <- norm_eucl(m)
```

```
cl<- kmeans(m_norm, 21)
```

```
table(cl$cluster)
```

```
plot(prcomp(m_norm)$x, col=cl$cl)
```

```
cl$withinss
```

```
cl$tot.withinss
```



Results for 20 Groups data :

Number of clusters is 2, based on NBClust

```
nc <- NbClust(dtm, min.nc=2, max.nc=30, method="kmeans")
```

```
table(nc$Best.n[1,])
```

0	2	3	5	7	12	17	20	21	22	30
2	6	5	1	2	1	1	1	4	1	2




Plotting data (Term Frequency – Inverse Document Frequency):

```
table(c1$cluster)
```

1	2
987	2013

SSE : 2941.527  
There exists overlapping across clusters



## LSA : Plotting document vector from SVD :


D = 50 (Reducing the dimension to 50 and considering only highest 50 Eigen values in diagonal vector D)

```
>table(cl$cluster)
```

```
 1  2  
997 2003
```

SSE : 2062.121

There c lusters are separated out




D = 100 (Reducing the dimension to 100 and considering only highest 100 Eigen values in diagonal vector D)

```
>> table(cl$cluster)
```

```
 1  2  
2052 948
```

SSE : 2470.836

There is some overlapping among the two clusters, but better than tf-idf.




D = 200(Reducing the dimension to 200 and considering only highest 20 Eigen values in diagonal vector D)

```
> table(cl$cluster)
```

```
 1  2
947 2053
```

SSE : 2696.556

There is some overlapping among the two clusters, but better than tf-idf.




LDA (Plotting the topic probabilities) :

K=2

```
> table(cl$cluster)
```

1	2
1216	1784

SSE : 1933.129




K=3

```
table(cl$cluster)
```

1	2	3
962	1694	344

SSE : 1831.981




K=21

```
> table(cl$cluster)
```

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
20	21																	
453	175	23	94	57	109	950	104	196	33	225	32	11	57	8	81	15	255	23
52	47																	

SSE : 1224.862



```

> lda.terms <- as.matrix(terms(lda,))
> head(lda.terms,20)#top terms in the topics learnt
      [,1]
Topic 1  "benifit"
Topic 2  "broadcast"
Topic 3  "dri"
Topic 4  "ottawa"
Topic 5  "ether"
Topic 6  "jon"
Topic 7  "game"
Topic 8  "aplandmalabcca"
Topic 9  "sat"
Topic 10 "clutch"
Topic 11 "xnewsread"
Topic 12 "gonna"
Topic 13 "coegalonlarcnasagov"
Topic 14 "bondracot"
Topic 15 "pgfsrlcacsusledu"
Topic 16 "ulf"
Topic 17 "figur"
Topic 18 "gray"
Topic 19 "german"
Topic 20 "perkyacsbuedu"

```

## Comparison and Results for 20NG datasets :

We can clearly see that the SSE of the LDA is lower than that of LSA and tf-idf

Model		SSE
Tf-idf		2941.527
LSA	D = 50	2062.121
	D = 100	2470.836
	D = 200	2696.556
LDA	K = 2	1933.129



Worcloud :




Number of clusters is set to 4.

### Plotting data (Term Frequency – Inverse Document Frequency):

```
> table(c1$cluster)
```

1	2	3	4
736	547	31	46

SSE : 3654.87




LSA :

D = 50

SSE: 2201.57

Most Representative words :

"war"	"und"	"laden"	"den"	"bin"	"mein"	"a
mbient"	"gym"	"buffet"	"asia"			




D = 100

SSE : 2954.02

Most Representative words :


"war" "und" "laden" "den" "bin" "mein" "ambient" "gym"  
"buffet" "asia"



D = 200


SSE : 3304.342

Most Representative words : "war" "und" "laden" "den" "b  
in" "mein" "ambient" "gym" "buffet" "asia"



LDA(Plotting the topic probabilities) :

SSE : 1869.126



Comparison and Results for 20NG datasets :

We can clearly see that the SSE of the LDA is lower than that of LSA and tf-idf

Model		SSE
Tf-idf		3654.87
LSA	D = 50	2201.57
	D = 100	2959.02
	D = 200	3304.342
LDA	K = 4	1869.126

## Results Summary:

Data	Model	SSE
20NewsGroups	Tf-idf	2941.527
	LSA D = 50	2062.121
	D = 100	2470.836
	D = 200	2696.556
	LDA K = 2	1933.129
Yelp	Tf-idf	3654.87
	LSA D = 50	2201.57
	D = 100	2959.02
	D = 200	3304.342
	LDA K = 4	1869.126

## Conclusion:

I can conclude from the experiments above that LDA perform better than LSA which performs better than TF-IDF.

TF-IDF just uses the term frequencies and the inverse document frequencies of terms , so unrelated words or documents can fall into the same category just because they have similar frequencies across the corpus.

LSA is an effective technique that brings down the hidden information be plotting the data into lower dimensions as this process DE noises the data and makes it more cleaner and meaningful.

LDA takes few steps ahead and assign probabilities to terms and topics such that there are chances of a term belonging to various topics. Similarly to the topics in the model This provides more flexibility to the model.

I understood that LDA is a better model than LSA. I also learnt to implement these techniques in R. It was a good practice, learning and understanding process.